

TP3: Introduction to Software Logging and Observability

Exercise 1: Logging a (*backend*) Java Application for User Profiling

Question 1

Create a simple backend application that allows you to :

1. create a user with an ID, name, age, email, and password.
2. provide a user with a menu through which (s)he can :
 - display products in a repository, where every product has an ID, a name, a price, and a expiration date.
 - fetch a product by its ID (if no product with the provided ID exists, an exception must be thrown).
 - add a new product (if a product with the same ID already exists, an exception must be thrown).
 - delete a product by its ID (if no product with the provided ID exists, an exception must be thrown).
 - update a product's info (if no product with the provided ID exists, an exception must be thrown).

The backend can be consumed with a simple frontend CLI. Also, adding a database (e.g., MongoDB) is a bonus, which would be very beneficial for the second bonus exercise.

Question 2

1. watch the following videos :
 - [Logging in Java](#)
 - [Logback vs SLF4J vs Log4J2 - what is the difference? Java Brains Brain Bytes](#)
 - [How to do logging in Spring Boot - Brain Bytes](#)
2. choose a logging utility among the ones seen in the videos and play around with it a bit to understand its different facets. **Note:** The usage of Log4j2 or SLF4J is strongly recommended.

Question 3

Use Eclipse JDT or Spoon (*recommended*) to log the code of your backend application using the logging utility chosen in the last question, such that the generated logs can be leveraged to create profiles of users as follows :

- a profile for those that mostly performed read operations on your database.
- a profile for those that mostly performed write operations on your database.
- a profile for those that searched for the most expensive products in your database.

The definition of your profile's structure is up to you. The most important information to include is the user's info, and the marking features of your profile (e.g., *the read/write operations performed by a user for the mostly performed read/write profiles*). The storage format for your profiles and their associated data is also up to you.

Note: The usage of JSON is strongly recommended due to its lightweight qualities and ease of serialization/deserialization to/from some persistence entity (e.g., file system, database, etc.).

Question 4

Define and execute a sequence of execution scenarios with different users to generate your logs. For example, you can create 10 users, and let each user execute around 20 different scenarios involving the above operations with different input values. Make sure that your scenarios are diverse enough to simulate a real-world experience to have properly crafted profiles at the end.

Question 5

Propose a way to parse the generated logs and extract the required information to construct your user profiles. Note that structured logs require the reification of an LPS and the definition of its construction and printing mechanisms (e.g., `Timestamp: <date_time>`, `Event: <event_info>+`, `User: <user_info>+`, `Action: <method_info>+`, etc.). In this case, each part of the LPS' content can be built separately and then aggregated to form the resulting LPS (*think about the **Builder** Design Pattern*). You can also choose to store and display them as JSON files if you wish.

Exercise 2: Tracing a (*frontend*) web application with OpenTelemetry for Service Visualization (*Bonus*)

Question 1

Create a simple web interface for your backend application (e.g., Angular, React, etc.). Think about exposing a REST API from your backend application to be consumed by your front-end application.

Question 2

Use a [frontend OpenTelemetry instrumentation agent](#) to (*automatically and/or manually*) **trace** the code of your frontend application, such that the generated logs can be visualized in a backend of your choice (e.g., Zipkin, ELK Stack). The generated traces should include information about the frontend operations leading up to a transaction that gets executed in the backend.

Question 3

Define and execute a sequence of execution scenarios with different users to generate your traces. Make sure that your scenarios are diverse enough to simulate a real-world experience.

Question 4 (*For the brave ones*)

1. Add a [backend OpenTelemetry instrumentation agent](#) to (*automatically and/or manually*) **trace** the code of your backend application, such that the generated traces can be visualized in a observability visualization backend of your choice (e.g., Zipkin). The generated traces should include information about the backend operations in response to front-end sent requests.
2. Try to find a way to connect the front-end and back-end traces and visualize the entire end-to-end traces of your application.